

INDEX

Sl.no.	Title	Sign
21/3/24 1.	Import & Export data using Pandas library	-10 ✓✓
28/3/24 2.	Various Data preprocessing techniques	-10 ✓✓
4/4/24 3.	Simple & multiple linear regression	-10 ✓✓
18/4/24 4.	TDS	-10 ✓✓
25/4/24 5.	Logistic regression	-10 ✓✓
9/5/24 6.	ANN - SVM	-10 ✓✓
23/5/24 7.	ANN, Randomforest	-10 ✓✓
23/5/24 8.	Randomforest, AdaBoost	-16 ✓✓
30/5/24 9.	Clustering kmeans	-10 ✓✓
30/5/24 10.	Dim Red → PCA	-10 ✓✓

21/3/24

Labs 1

To import & export data using Pandas library functions.

(i) importing data

import pandas as pd

```
airbnb_data = pd.read_csv("data/listings-austin.csv")  
airbnb_data.head()
```

(ii) Reading data from URL

```
url = "https://archive.ics.uci.edu/ml/machine-learning-  
database/iris/iris.data"
```

```
col_names = ["sepal-length-in-cm", "sepal-width-in-cm",  
             "petal-length-in-cm", "petal-width-in-cm",  
             "class"]
```

```
iris_data = pd.read_csv(url, names=col_names)
```

```
iris_data.head()
```

(iii) Exporting dataframe to CSV file

```
iris_data.to_csv("cleaned_iris_data.csv")
```

OUTPUT:

	sepal-length-in-cm	sepal-width-in-cm	petal-length-in-cm	petal-width-in-cm	class
0	5.1	3.5	1.4	0.2	Iris-setosa
1					
2					
3					
4					

Waz
21/3/24

28/3/24 End-to-End Machine Learning Project

- 1. Look at the big picture
- 2. Get the Data
- (1) Download the data

```
import os
import tarfile
import urllib
```

```
DOWNLOAD_ROOT = "https://raw.githubusercontent.com/ml-pipeline/tutorials/master/datasets/housing"
housing_path = os.path.join("data", "1")
housing_url = download_root + "datasets/housing/housing.tgz"
def fetch_housing(housing_url=housing_url,
                  housing_path=housing_path):
    os.makedirs(name=housing_path, exist_ok=True)
    egg_path = os.path.join(housing_path,
                           "housing.tgz")
    fetch_housing_data()

```

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
housing.hist(bins=50, figsize=(20,15))
plt.show()
```



```
housing[["income_cat"]] = pd.cut(x=housing['median_income'],
                                 bins=[0, 1.5, 3.45, 6, np.inf],
                                 labels=[1, 2, 3, 4, 5])
housing[["income_cat"]].hist()
```

3. Discover & Visualize the data to gain insights

```
strat_train_set, strat_test_set = stratified_set_index().loc['train'].copy(),
strat_test_set.reset_index(inplace=True)
housing = strat_train_set.copy()
housing.shape
```

```
housing.plot(kind='scatter', x='longitude', y='latitude')
plt.show()
corr_matrix = housing.corr()
corr_matrix['median_house_value'].sort_values(ascending=False)
```

```
from pandas.plotting import scatter_matrix
attributes = ['median_house_value', 'median_income',
              'total_income', 'housing_median_age']
scatter_matrix(frame=housing[attributes],
               figsize=(12, 8))
plt.show()
```

~~project-27-end-to-end.ipynb - step 3~~

4. Prepare the Data for Machine Learning Algorithms

```
from sklearn.impute import SimpleImputer  
imputer = SimpleImputer(strategy='median')
```

```
housing_num = housing.drop(['ocean-proximity'], axis=1)  
imputer.fit(housing_num)
```

imputer.statistics -

```
housing_num.median().values
```

```
X = imputer.transform(housing_num)
```

X.shape

```
housing_tr = pd.DataFrame(data=X, index=housing.index,  
                           columns=housing_num.columns)
```

housing_tr.head()

```
housing_cat = housing[['ocean-proximity']]
```

housing_cat.head()

```
from sklearn.pipeline import Pipeline
```

```
from sklearn.preprocessing import StandardScaler
```

```
num_pipeline = Pipeline([
```

```
(imputer, SimpleImputer(strategy='median'))
```

```
('atributes-adder', CombinedAttributes
```

```
Adder())
```

```
('std-scalar', StandardScaler())
```

])

housing_num_tr = num_pipeline.fit_transform(housing_num)

housing_num_tr.shape

(500, 32)

WV324
74

5. select & train a model:
- from sklearn.linear_model import LinearRegression
lin-reg = LinearRegression()
 - lin-reg.fit(x= housing_prepared, y=housing_labels)
 - housing_predictions = lin-reg.predict(housing_prepared)
 - lin_mse = mean_squared_error(housing_labels, housing_predictions)
 - lin_rmse = np.sqrt(lin_mse)
 - lin_rmse
 - tree-reg = DecisionTreeRegressor()
 - tree-reg.fit(x=housing_prepared, y=housing_labels)
 - forest-reg = RandomForestRegressor()
 - forest-reg-fit(x=housing_prepared, y=housing_labels)

6. Fine-tune Your Model
- grid-search.best_params_
 - grid-search.best_estimator_
 - cat-encoder = full_pipeline.named_transformers_
 - final-model = grid-search.best_estimator_[^{'cat'}]
 - X-best-prepared = full_pipeline.transform($\lambda = \text{X-test}$)
 - final_rmse = np.sqrt(final_mse)
 - final_rmse
 - squared_errors = ($y_{\text{test}} - \text{final-predictions}$) ** 2

7. Launch, monitor & maintain your system.

4/14/24

Week 3

SLR

```
df-sal = pd.read_csv('Salary-Data.csv')
df-sal.head()
```

```
df-sal.describe
```

```
plt.title('Salary Dist Plot')
sns.distplot(df-sal['Salary'])
plt.show()
```

```
plt.scatter(df-sal['Years Exp'], df-
```

```
color = 'lightcoral'
```

```
plt.title('Sal vs Exp')
```

```
plt.ylabel('Salary')
```

```
plt.xlabel('Years of Exp')
```

```
plt.box(False)
```

```
plt.show()
```

```
x = df-sal.iloc[:, :1]
```

```
y = df-sal.iloc[
```

```
X-train, X-test, Y-train, Y-test =
```

```
train-test = split(x, y, test-size = 0.2, random-state = 0)
```

```
regression = Linear Regression()
```

```
regression.fit(X-train, Y-train)
```

```
plt.scatter(X-train, Y-train, color = 'lightcoral')
```

```
plt.plot(X-train, Y-pred-train)
```

```
print(f'coff: {regression.coef_}')
```

~~```
print(f'intercept: {regression.intercept_}')
```~~

## multiple linear regression

df\_start = pd.read\_csv('')  
df\_start.head()

df\_start.describe()

plt.title('Profile Disk Plot')

plt.show

plt.scatter(df\_start['R9D'], df\_start['R9D'])

df\_start['Profilt'].color = 'light coral'

plt.box(False)

plt.show()

x = df\_start.iloc[:, :-1].values

y = df\_start.iloc[:, -1].values

regression = LinearRegression()

regressor.fit(x-train, y-train)

y-pred = regressor.predict(x-test)

result = np.concatenate((y-pred.reshape

(len(y-pred), 1), y-test.reshape(len(y-test))))

print(result)

18/4/24

Week 4

ID3

```
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier.
plot tree
import matplotlib.pyplot as plt
```

```
iris = load_iris()
```

```
X = iris.data
```

```
y = iris.target
```

```
feature_names = iris.feature_names
```

```
target_names = iris.target_names
```

```
X_train, X_test, Y_train, Y_test =
```

```
train_test_split(X, Y, test_size=0.4,
random_state=42)
```

```
clf = DecisionTreeClassifier(criterion='gini').
```

```
accuracy = clf.score(X_test, Y_test)
```

```
print("Accuracy", accuracy)
```

```
plt.figure(figsize=(12, 8))
```

```
plt.show()
```

Q.P:

Accuracy = 0.983

W/Kin

25/4/24

Weeks

### logistic regression

```
import pandas as pd
```

```
from matplotlib import pyplot as plt
```

```
%matplotlib inline
```

```
df = pd.read_csv("insurance-data.csv")
```

```
df.head()
```

```
df['age'].scatter(df['age'], df['bought_insurance'], marker='+', color='red')
```

```
from sklearn.model_selection import train_test_split
```

```
Perform the split
```

```
X_train, X_test, Y_train, Y_test =
```

```
train_test_split(df[['age']], df['bought_insurance'], train_size=0.8)
```

```
from sklearn.linear_model import LogisticRegression
```

```
model = LogisticRegression()
```

```
model.fit(X_train, Y_train)
```

```
y_predicted = model.predict(X_test)
```

```
model.predict_proba(X_test)
```

```
model.score(X_test, Y_test)
```

Output:

Accuracy : 0.83259147

Analyze through Linear Regression

model-coef - indicates value of  $m$  in  $y = mx + b$  eq.  
model-coef -  $\text{coef}(\text{slope}) = [0.14544708]$

model-intercept - indicates value of  $b$  in  $y = mx + b$   
model-intercept -  $[ -5.44733781 ]$

Analyze through Sigmoid function:

import math

def sigmoid(x):

return  $1 / (1 + \text{math.exp}(-x))$

def prediction\_function(age):

$z = 0.042 * \text{age} - 1.53 + 0.04150133 \approx$   
 $0.042 \text{ and } -1.52726963 \approx -1.53$

$y = \text{sigmoid}(z)$

return y

age = 35

prediction\_function(age)

prob: 0.4850044983805899

age = 43

prediction\_function(age)

prob: 0.568565299077705

QWV  
25-4-2023

7/15/23

KNN classifier:

- ① import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
plt.style.use('ggplot')
- ② df = pd.read\_csv('/diabetes.csv')  
df.head  
df.shape
- ③ X = df.drop(['Outcome'], axis=1).values  
Y = df['Outcome'].values
- ④ from sklearn.model\_selection import train\_test\_split  
X\_train, X\_test, y\_train, y\_test =  
train\_test\_split(X, y, test\_size=0.4,  
random\_state=42,  
stratify=Y)
- ⑤ from sklearn.neighbors import KNeighborsClassifier  
neighbors = np.arange(1, 9)  
train\_accuracy = np.empty(len(neighbors))  
test\_accuracy = np.empty(len(neighbors))  
knn.fit(X\_train, y\_train)  
train\_accuracy[i] = knn.score(X\_train, y\_train)  
test\_accuracy[i] = knn.score(X\_test, y\_test)

(6) plt.plot(neighbors, train\_accuracy, label='Train')  
plt.plot(neighbors, test\_accuracy, label='test')  
plt.show()

(7) plt.scatter(neighbors, train\_accuracy, color='k')  
plt.show()

(8) knn.fit(x\_train, y\_train)

knn.score(x\_test, y\_test)

O/P Accuracy = 0.7348

WV  
95-21

SVM

from sklearn import datasets

cancer = datasets.load\_breast\_cancer()

# enter features & print

print(cancer.target)

from sklearn.model\_selection import train\_test\_split

X-train, X-test, y-train, y-test =

train\_test\_split(cancer.data, cancer.target,

test\_size=0.3, random\_state=109)

from  
import sklearn import SVM

clf = SVC(kernel='linear')

y-pred = clf.predict(X-test)

from sklearn import metrics

print(metrics.accuracy\_score(y-test, y-pred))

print(metrics.precision\_score(y-test, y-pred))

print(metrics.recall\_score(y-test, y-pred))

import seaborn as sns

sns.scatterplot(x=cancer.data[:, 0],

y=cancer.data[:, 1], hue=cancer.target)

clf

Accuracy: 0.9764

W M  
W M

20/5/20  
 Vectors  
 Artificial neural network (3 layers)  
 import numpy as np  
 $x = \text{np.array}((1, 2, 3, 4, 5, 6, 7, 8, 9), \text{shape}=(3, 3))$   
 $y = \text{np.array}(10, 11, 12, 13, 14, 15, 16, 17, 18), \text{shape}=(3, 3)$   
 $x = x / \text{np.max}(x, \text{axis}=0)$   
 $y = y / 100$   
 epoch = 5000  
 lr = 0.1  
 input\_neurons = 3  
 hiddenlayer\_neurons = 3  
 output\_neurons = 1  
 $wh = \text{np.random.uniform}(\text{size}=(\text{input-layer-neurons},$   
 $\text{hidden-layer-neurons}))$   
 $bh = \text{np.random.uniform}(\text{size}=(1, \text{hidden-layer-neurons}))$

```

def sigmoid(x):
 return 1 / (1 + np.exp(-x))

def derivatives_sigmoid(x):
 return x * (1 - x)

for i in range(epoch):
 hinp1 = np.dot(x, wh)
 hinp = hinp1 + bh
 hlayer_act = sigmoid(hinp)
 output = sigmoid(hinp)
 EO = y - output
 out_grad = derivatives_sigmoid(output)
 wh -= EO * out_grad * x
 bh -= EO * out_grad

```

$Eh = d \cdot \text{output} \cdot (1 - \text{output})$

hidden grad = derivative - sigmoid ( $\text{H}_j$   $\text{grad}$  -  $\sigma$ )

$\Delta$ -hiddenlayer =  $E_{HP} \times$  hiddengrad

$nh += X \cdot T \cdot \text{dot}(\Delta\text{-hiddenlayer}) \times lr$

// print  $I/P \rightarrow O/P$

O/P : Input

$\begin{bmatrix} [0.66667 \ 1.] \\ [0.33333 \ 0.55556] \\ [1. \ 0.66667] \end{bmatrix}$

Actual Output:

$\begin{bmatrix} [0.92] \\ [0.86] \\ [0.87] \end{bmatrix}$

Predicted Output

$\begin{bmatrix} [0.8196977] \\ [0.7969237] \\ [0.81616067] \end{bmatrix}$

QW  
23-5-24

## 9a) Random Forest - Prog 7

// import libraries

data = pd.read\_csv('Iris.csv')

X = data.drop(columns=['Species'])

y = data['Species']

X-train, X-test, y-train, y-test =

train-test-split(X, y, test\_size=0.2, random\_state=42)

rf\_classifier = RandomForestClassifier(n\_estimators=100, random\_state=42)

rf\_classifier.fit(X-train, y-train)

y-pred = rf\_classifier.predict(X-test)

accuracy = accuracy\_score(y-test, y-pred)

// print accuracy, classification report

feature-importances = rf\_classifier.feature\_importances\_

features = X.columns

importance-df = pd.DataFrame({'Feature': features,

'Importance': feature-importances}).

sort-values(by='Importance')

ascending=False

print(importance-df)

Output: Accuracy : 100.00%.

WU  
23-5-21

## q6) AdaBoost

```
// import necessary libraries
data = pd.read_csv('content/iris.csv')
X = data.drop(columns = ["species"])
y = data["species"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2)
random_state = 42

adaBoostclf = AdaBoostClassifier(n_estimators = 150, learning_rate = 1)

adaBoostclf.fit(X_train, y_train)

y_pred = adaBoostclf.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
```

// print Acc

Op:

Accuracy : 1.0

W/SU  
23/5/21

20/10/24

Week 8

LAB PROJ 10

K-Means Algorithm to cluster a set of data stored in a .csv file

import mlxtend.lib as plt

from sklearn import datasets

from sklearn.cluster import KMeans

import pandas as pd

import numpy as np

iris = datasets.load\_iris()

X = pd.DataFrame(iris.data)

X.columns = ['Sepal-length', 'Sepal-width',  
'Petal-length', 'Petal-width']

y = pd.DataFrame(iris.target)

y.columns = ['Targets']

model = KMeans(n\_clusters=3)

model.fit(X)

plt.figure(figsize=(10, 10))

color\_map = np.array(['red', 'lime', 'black'])

plt.subplot(2, 2, 1)

plt.scatter(X.Petal-Length, X.Petal-width)

c = color\_map[y.Targets], s=

plt.title('Real Clusters')

plt.xlabel('Petal length')

plt.ylabel('Petal width')

plt.subplot(2, 2, 2)

plt.scatter(X.Petal-Length, X.Petal-width)

c = color\_map[model.labels], s=

plt.title('K-means clustering')

plt.xlabel('Petal length')

plt.ylabel('Petal width')

3/5/24

Week 8

LAB 11 Dimensionality reduction using PCA

import matplotlib.pyplot as plt

import pandas as pd

import seaborn as sns

%matplotlib inline

from sklearn.datasets import load\_breast\_cancer

cancer = load\_breast\_cancer()

cancer.keys()

print(cancer['DESCR'])

df = pd.DataFrame(cancer['data'], columns=cancer['feature names'])

from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

scaler.fit(df)

scaled\_data = scaler.transform(df)

from sklearn.decomposition import PCA

pca = PCA(n\_components=2)

pca.fit(scaled\_data)

pca.components\_.shape

x\_pca = pca.transform(scaled\_data)

scaled\_data.shape

x\_pca.shape

plt.figure(figsize=(6, 6))

plt.scatter(x\_pca[:, 0], x\_pca[:, 1], c=cancer['target'])

plt.xlabel('First principal component')  
cmnp = 'principal'

plt.ylabel('Second PC')

3/5/24